

Administrivia

- Canceling my office hours this week
- Phil and I both have to be out of town
- Sachin Katti will give guest lecture on Coding Thursday

Congestion Control Revisited

- **Congestion is when the input rate \gg output rate**
 - In TCP, flow control window ensures sender does not exceed rate at which receiver consumes data
 - What if senders exceed a router's maximum output rate?
- **What should routers do? Make sender slow down**
- **TCP sending rate = window-size/RTT, so 2 options:**
 1. Increase RTT – buffer more packets \Rightarrow more queuing delay
 2. Reduce window size – happens if router drops packets
- **Recall TCP reacts to packet loss by shrinking congestion window**
 - Triple duplicate ack: halve window, enter CA state
 - Timeout: set window to 1, enter SS state


Router design issues

- **Scheduling discipline**
 - Which of multiple packets should you send next?
 - May want to achieve some notion of fairness
 - May want some packets to have priority
- **Drop policy**
 - When should you discard a packet?
 - Which packet to discard?
 - Some packets more important (perhaps BGP)
 - Some packets useless w/o others (IP fragments)
- **Need to balance throughput & delay**
 - Could minimize/eliminate drops with enormous buffers
 - But queuing delay highly frowned upon (interactive apps)

Overview

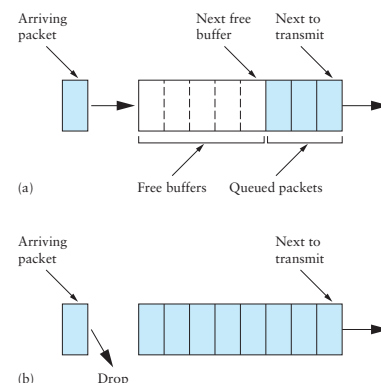
- **How routers queue affects how TCP and other protocols behave**
- **Two router questions: drop policy, scheduling policy**
- **Reducing congestion through content distribution**
 - Clients can cache
 - Services can use a CDN

Congestion at Router



- **Router goals**
 - Prioritize who gets limited resources
 - Somehow interact well with TCP

Example: FIFO tail drop




- **Differentiates packets only by when they arrive**
 - Packet dropped if queue full when it arrives

Tail drop issues

- **When stable, queue will always be nearly full**
 - Guarantees high latency for all traffic
- **Possibly unfair for flows with small windows**
 - E.g., small flow (< 4 packages) may be stuck in backoff, while larger flows can use fast retransmit to recover
- **Window synchronization**
 - Consider many flows in a stable configuration
 - New flow comes in, causes a bunch of packet losses
 - Existing flows all cut their windows together (underutilizing link)
 - Flows all grow their windows together until link again overloaded and many packets lost. Repeat...


Connectionless flows



- **Even in Internet, routers can have a notion of flows**
 - E.g., base on IP addresses & TCP ports (or hash of those)
 - *Soft state*—doesn't have to be correct
 - But if often correct, can use to form router policies


Scheduling Policy: Fair Queuing (FQ)

- Explicitly segregates traffic based on flows
- Ensures no flow consumes more than its share
 - Variation: weighted fair queuing (WFQ)
- Note: if all packets were same length, would be easy




What to optimize for?

- *Fairness* (in two slides)
- *High throughput* – queue should never be empty
- *Low delay* – so want short queues
- **Crude combination: $\text{power} = \text{Throughput}/\text{Delay}$**
 - Want to convince hosts to offer optimal load



Fairness



- **What is fair in this situation?**
 - Each flow gets 1/2 link b/w? Long flow gets less?
- **Usually fair means equal**
 - For flow bandwidths (x_1, \dots, x_n) , fairness index:

$$f(x_1, \dots, x_n) = \frac{(\sum_{i=1}^n x_i)^2}{n \sum_{i=1}^n x_i^2}$$

- If all x_i s are equal, fairness is one
- Weighted fairness is a simple extension
- **So what policy should routers follow?**

Fair Queueing Basics


- Keep track of how much time each flow has used link
- Compute how long a flow will have used link if it transmits next packet
- Send packet from flow which will have lowest use if it transmits
 - Why not flow with smallest use so far?
 - Because next packet may be huge (examples coming)

FQ Algorithm

- Suppose clock ticks each time a bit is transmitted
- P_i : length of packet i
- S_i : time when packet i started transmission
- F_i : time when packet i finished transmission
- $F_i = S_i + P_i$
- When does router start transmitting packet i ?
 - If arrived before router finished packet $i - 1$ from this flow, then immediately after last bit of $i - 1$ (F_{i-1})
 - If no current packets for this flow, then start transmitting when arrives (call this A_i)
- Thus: $F_i = \max(F_{i-1}, A_i) + P_i$

FQ Algorithm (cont)

- For multiple flows
 - Calculate F_i for each packet that arrives on each flow
 - Treat all F_i s as timestamps
 - Next packet to transmit is one with lowest timestamp
- Not perfect: can't preempt current packet
- Example:



FQ Algorithm (cont)

- One complication: inactive flows are penalized ($A_i > F_{i-1}$)
- Over what interval do you consider fairness?
 - Standard algorithm considers no history
 - Each flow gets fair share while packets queued
- Solution: $B_i = P_i + \max(F_{i-1}, A_i - \delta)$
- δ = interval of history to consider

Fair Queueing Importance


- "Our packet-by-packet transmission algorithm is simply defined by the rule that, whenever a packet finishes transmission, the next packet is the one with the smallest F_i^α ."
- But, fair queueing not used in core routers: finding min F in hundreds of thousands of flows is expensive. Can be used on edge routers and low speed links.

Drop Policy: Random Early Detection (RED)


- Notification of congestion is implicit in Internet
 - Just drop the packet (TCP will timeout)
 - Could make explicit by marking the packet (ECN extension to IP allows routers to mark packets)
- Early random drop
 - Don't wait for full queue to drop packet
 - Instead, drop packets with some *drop probability* whenever the queue length exceeds some *drop level*
 - Prevents global window synchronization: many TCP flows speed up, all have packets dropped, all slow down, etc.

RED Details

- Compute average queue length
 - $\text{AvgLen} = (1 - \text{Weight}) \cdot \text{AvgLen} + \text{Weight} \cdot \text{SampleLen}$
 - $0 < \text{Weight} < 1$ (usually 0.002)
- SampleLen is queue length each time a packet arrives



AvgLen



- Smooths out AvgLen over time
 - Don't want to react to instantaneous fluctuations


RED Details (cont)

- Two queue length thresholds:
 - if AvgLen ≤ MinThreshold then enqueue the packet
 - if MinThreshold < AvgLen < MaxThreshold then calculate probability P drop arriving packet with probability P
 - if MaxThreshold ≤ AvgLen then drop arriving packet

RED Details (cont)

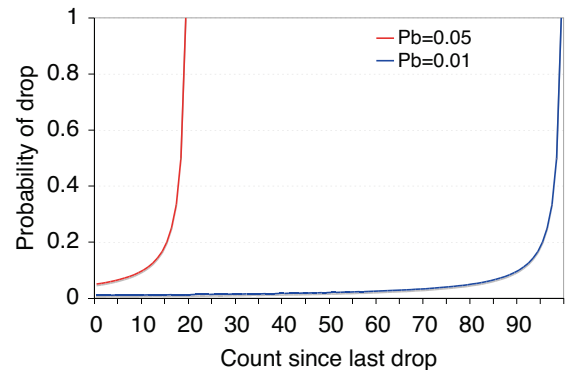
- Computing probability P

$$P_b = \text{MaxP} \cdot \frac{(\text{AvgLen} - \text{MinThreshold})}{(\text{MaxThreshold} - \text{MinThreshold})}$$



- Actual drop probability based on time since last drop
 - count = # pkts since drop or MinThreshold < AvgLen < MaxThreshold
 - $P = P_b / (1 - \text{count} \cdot P_b)$
 - Space out drops, separate when to drop from which to drop

What P looks like



Tuning RED

- Probability of dropping a particular flow's packet(s) is roughly proportional to the share of the bandwidth that flow is currently getting
- MaxP is typically set to 0.02
- If traffic is bursty, then MinThreshold should be sufficiently large to allow link utilization to be maintained at an acceptably high level
- Difference between two thresholds should be larger than the typical increase in the calculated average queue length in one RTT; setting MaxThreshold to twice MinThreshold is reasonable for traffic on today's Internet

Queueing Today

- Cisco IOS
 - Scheduling: FIFO, FQ, WFQ, Custom queueing (patterns)
 - Drop policy: Tail drop, weighted random early detection

2-minute stretch



Content distribution

- **How can end nodes reduce load on bottleneck links?**
 - Congestion makes net slower – nobody wants this
- **Client side**
 - Many people from Stanford might access same web page
 - Redundant downloads a bad use of Stanford's net connection
 - Save resources by caching a copy locally
- **Server side**
 - Not all clients use caches
 - Can't upload unlimited copies of same data from same server
 - Push data out to content distribution network

Examples

- **Web browser caches recently accessed objects**
 - E.g., allows "back" button to operate more efficiently
- **Web proxies cache recently accessed URLs**
 - Save bandwidth/time when multiple people locally access same remote URL
- **DNS resolvers cache resource records**
- **Network file systems cache read/written data**
- **PDA caches calendar stored in Desktop machine**

Caching

- Many network apps. involve transferring data
- **Goal of caching: Avoid transferring data**
 - Store copies of remotely fetched data in *caches*
 - Avoid re-receiving data you already have
- **Caching concerns keeping copies of *data***

One approach: TTLs

- **Eventual consistency**
- **Source controls how long data can be cached**
 - Can adjust trade-off: Performance vs. Consistency
- **Example: TTLs in DNS records**
 - When looking up `vine.best.stanford.edu`
 - CNAME record for `vine.best.stanford.edu` has very short TTL—value frequently updated to reflect load averages & availability
 - NS records for `best.stanford.edu` has long TTL (can't change quickly, and `stanford.edu` name servers want low load)
- **Example: HTTP reply can include Expires: field**

Cache consistency

- **Problem: What happens when objects change?**
- **Is cached copy of data is up to date?**
- **Stale data can cause problems**
 - E.g., don't see edits over a network file system
 - Get wrong address for DNS hostname
 - Shopping cart doesn't contain new items on web store
- **Can have various degrees of consistency**

Polling

- **Check with server before using a cached copy**
 - Check requires far less bandwidth than downloading object
- **How to know if cache is up to date?**
 - Objects can include version numbers
 - Or compare time-last-modified of server & cached copies
- **Example: HTTP If-Modified-Since: request**
- **Sun network file system (NFS)**
 - Caches file data and attributes
 - To validate data, fetch attributes & compare to cached

Leases

- **Leases – promise of callback w. expiration time**
 - E.g., Download cached copy of file
 - Server says, “For 2 minutes, I’ll let you know if file changes”
 - Or, “You can write file for 2 minutes, I’ll tell you if someone reads”
 - Client can renew lease as necessary
- **What happens if client crashes or network down?**
 - Server might need to invalidate client’s cache for update
 - Or might need to tell client to flush dirty file for read
 - Worst case scenario – only need to wait 2 minutes to repair
- **What happens if server crashes?**
 - No need to write leases to disk, if rebooting takes 2 minutes
- **Used by Google’s internal naming/lock service (Chubby)**
- **Gray, Cheriton won test of time award for leases work done here at Stanford**

Why CDNs succeed more (compared to web caches)

- **Incentives**
- **Content provider (e.g., Microsoft) uses/deploys CDN: wants to improve performance and reduce costs**
- **End user (e.g., network administrator) uses/deploys cache: wants to reduce external traffic**

Callbacks

- **Polling may cause scalability bottleneck**
 - Server must respond to many unnecessary poll requests
- **Example: AFS file system stores software packages**
 - Many workstations at university access software on AFS
 - Large, on-disk client caches store copies of software
 - Binary files rarely change
 - Early versions of AFS overloaded server with polling
- **Solution: Server tracks which clients cache which files**
 - Sends *callback* message to each client when data changes

Content Distribution Network (CDN)

- **Network of computers that replicate content across the Internet**
- **Bringing content closer to requests can improve performance**
- **All users communicate with Redmond to download Microsoft SP**
 - Bottleneck: pipes to Redmond
- **Microsoft pushes SP to many hosts around the country**
 - Uses only local (not shared) capacity
- **Actively pushes data into the network**

Akamai

- **Challenge: static host name needs to point to different servers based on location**
- **Akamai servers cache content (images, videos, etc.)**
- **Uses DNS to direct clients to “close” servers**
- **Specifically, points clients to close NS servers**
- **Different NS servers provide different host lookups**

Caches and load balancing

- **Let's say you are Akamai**
 - Clusters of server machines running web caches
 - Caching data from many customers
 - Proxy fetches data from customer's *origin server* first time it gets request for a URL
- **Chose cluster based on client network location**
- **How to choose server within a cluster?**
- **Don't want to chose based on client... low hit rate**
 - N servers in cluster means N cache misses per URL
- **Also don't assume proxy servers 100% reliable**

Straw man: Modulo hashing

- **Say you have N proxy servers**
- **Map requests to proxies as follows:**
 - Number servers from 1 to N
 - For URL `http://www.server.com/web_page.html`, compute $h \leftarrow \text{HASH}(\text{"www.server.com"})$
 - Redirect clients to proxy $\# p = h \bmod N$
- **Keep track of load on each proxy**
 - If load on proxy $\# p$ is too high, with some probability try again with different hash function
- **Problem: Most caches will be useless if you add/remove proxies, change value of N**

Consistent hashing [Karger]

- **Use circular ID space based on circle**
 - Consider numbers from 0 to $2^{160} - 1$ to be points on a circle
- **Use circle to map URLs to proxies:**
 - Map each proxy to several randomly-chosen points
 - Map each URL to a point on circle (hash to 160-bit value)
 - To map URL to proxy, just find successor proxy along circle
- **Handles addition/removal of servers much better**
 - E.g., for 100 proxies, adding/removing proxy only invalidates $\sim 1\%$ of cached objects
 - But when proxy overloaded, load spills to successors
 - When proxy leaves, extra misses disproportionately affect successors, but will be split among multiple successors
- **Can also handle servers with different capacities**
 - Give bigger proxies more random points on circle

Cache Array Routing Protocol (CARP)

- **Different URL \rightarrow proxy mapping strategy**
 - Let list of proxy addresses be p_1, p_2, \dots, p_n
 - For URL u , compute:
 $h_1 \leftarrow \text{HASH}(p_1, u), h_2 \leftarrow \text{HASH}(p_2, u), \dots$
 - Sort h_1, \dots, h_n . If h_i is minimum, route request to p_i .
 - If h_i overloaded, spill over to proxy w. next smallest h
- **Advantages over consistent hashing**
 - Spreads load more evenly when server is overloaded, if overload is just unfortunate coincidence
 - Spreads additional load more evenly when a proxy dies

Overview

- **How routers handle overload affects how TCP (and other protocols) behaves**
- **Two router questions: drop policy, scheduling policy**
- **Can reduce congestion through content distribution**
 - Clients can cache, need techniques for consistency
 - Services can use a CDN, load-balancing becomes important