

Lecture 10: TCP Friendliness, DCCP, NATs, and STUN

TCP Friendliness


Congestion Control

- TCP dynamically adapts its rate in response to congestion
- AIMD causes flows to converge to fair goodput
- But how do losses (e.g., bit errors) affect goodput?
- What about UDP?

Responding to Loss

- Set threshold to $\frac{cwnd}{2}$
- On timeout
 - Set cwnd to 1
 - Causes TCP to enter slow start
- On triple duplicate ACK (Reno)
 - Set cwnd to $\frac{cwnd}{2}$
 - Retransmit missing segment
 - Causes TCP to stay in congestion avoidance

Chiu Jain Phase Plots



Analyzing TCP Simply

- Assume all segments are MSS long
- Assume a packet loss rate p
- Assume a constant RTT
- Assume p is small (no timeouts)

Window Size

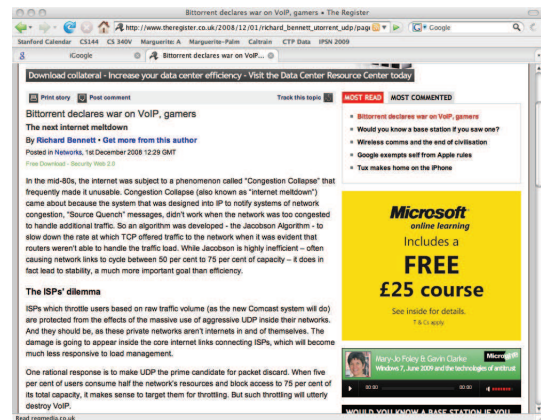
Analysis

- Window size W cuts to $\frac{W}{2}$ after a loss
- Grows to W after $\frac{W}{2}$ RTTs
- Goodput = $\frac{3}{4} \cdot W \cdot MTU \cdot \frac{1}{RTT}$

- $p = \frac{1}{(\frac{W}{2} + (\frac{W}{2} + 1) + \dots + W)}$
- $p \approx \frac{1}{\frac{3}{8}W^2}$
- $W \approx \sqrt{\frac{8}{3 \cdot p}}$
- Goodput = $\frac{3}{4} \cdot \sqrt{\frac{8}{3 \cdot p}} \cdot MTU \cdot \frac{1}{RTT}$
- Goodput = $\frac{1.22 \cdot MTU}{RTT \cdot \sqrt{p}}$
- Constant factor changes based on delayed acks, etc.

TCP Friendliness

- Don't want other protocols to disrupt TCP
- UDP happily shuts down TCP flows
- "TCP friendliness:" obeying TCP congestion control as per prior goodput equation
 - Does not imply acting like TCP
 - E.g., does not require abrupt window changes



ledbat WG


- "The LEDBAT WG is chartered to standardize a congestion control mechanism that should saturate the bottleneck, maintain low delay, and yield to standard TCP."
- TCP-friendliness is insufficient for modern P2P applications
 - Flow fairness, not application fairness
 - TCP fills queues
- Elastic workloads vs. inelastic workloads

DCCP

DCCP

- **Datagram Congestion Control Protocol (DCCP)** provides congestion control for unreliable datagrams (RFC 4340)
- **Connection-oriented protocol**
 - Request-response-ack establishment
 - Close-reset or CloseReq-Close-reset teardown
- **Counts packets, not bytes**

DCCP Segment



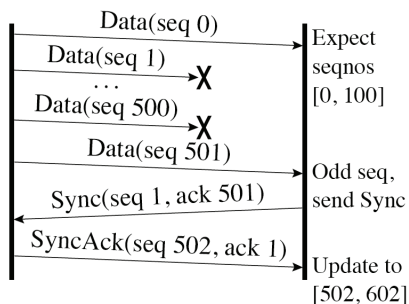
Sequence Numbers

- **Every DCCP packet uses a new sequence number**
 - Data
 - Acknowledgements
 - Control traffic
- **Acknowledgements are for *last packet received***
 - Not cumulative acknowledgements
 - Does not succinctly describe connection history
 - Options can give packet vectors


Synchronization

- **DCCP uses sequence number windows to protect from attacks**
- **Large bursts of losses cause packets to fall past windows**
- **Need to resynchronize**


Synchronization Exchange



Synchronization on Reset Problem



Congestion Control



- Defines Congestion Control IDs (CCIDs)
- Negotiated with change/confirm L/R options
- Each half-connection can have different congestion control
- CCID 2: TCP congestion control (AIMD) (RFC 4941)
- CCID 3: TCP-friendly congestion control (RFC 4942)

CCID 2

- **Uses TCP congestion control**
 - Maintains a *cwnd*, slow-start, etc.
- **Adds congestion control to acks**
 - Sender specifies an *AckRatio*, *R*
 - Ratio of data to ack packets (TCP with delayed ACKs is 2)
 - On detecting ack losses, double *R*
 - After $\frac{cwnd}{R^2 - R}$ lossless congestion windows, decrement *R*

CCID 3

- Uses TCP-friendly congestion control
- Uses a sending rate, rather than a congestion window
- Receiver sends feedback once per RTT, reporting loss rate
- If sender hears no feedback, halves sending rate
- Security issue with loss rate reporting: report *loss intervals*, rather than just a loss rate, verifiable with ECN nonces

DCCP Today

- Numerous implementations
- IETF Standards Track
- Well suited to VoIP, Internet Gaming, etc.
- **Sees very little use**

2-minute stretch




Network Address Translator


Motivations and Complications

- There are only 2^{32} IP addresses
- Firewalls for security
- Breaks end-to-end (node does not know its external IP)
- Node might not even know if it's behind a NAT
- NAT needs to be able to dynamically assign mappings

NAT Example, Step 1

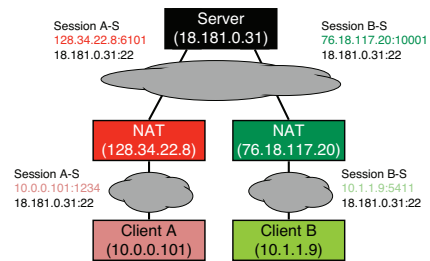


NAT Example, Step 2



NAT

• Network Address Translator




How a NAT Works

- Maps between global and local (IP,port) pairs
- Requires knowledge of transport packet format
- UDP datagram, TCP SYN
 - Can shut down TCP mapping on FIN+ACK
 - UDP requires timeouts (> 2 minutes, unless IANA says otherwise)
- RFC 4787/BCP 127 defines recommended behaviors


NAT Example, Step 3

NAT rewrites the source address and port so it seems it is the source




NAT Example, Step 4

NAT also creates a mapping so it can transform future packets correctly




NAT Example, Step 5

Server responds to SYN with SYN+ACK and NAT rewrites it




NAT Example, Step 6

Client B can also open a connection, and the NAT creates a mapping for it too



NAT Example, Step 7

In the case of TCP, NAT can discard the mapping on close or RST




Types of NAT (RFC 3849)

- Full Cone: no ingress filter (single local-external mapping)
- Restricted Cone: ingress filter on address
- Port Restricted: ingress filter on address/port
- Symmetric: different mappings for different external destinations
- Terminology is imperfect (static port mappings, etc.)

NAT Problems

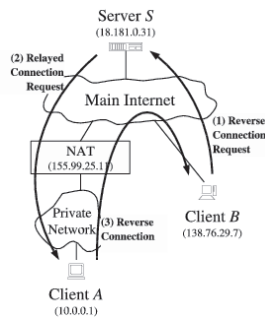
- Incoming TCP connections
- E.g., Skype



TCP Through NATs

- Server socket doesn't initiate traffic: NAT can't set up mapping
- Rendezvous servers (as in Skype)
- Connection reversal through rendezvous if only one is behind a NAT (rendezvous server asks un-NAT node to open a port so NAT node can connect)

TCP Reversal



More NAT Problems

- Port mapping: 0-1023 should map to 0-1023
- Port parity: even port → even port, odd port → odd port (RFC 3550: RTP uses even, RTCP uses odd)
- Hairpinning: packet from inside NAT can send packets to external (NAT) address and have portmap work

RFC 4787 and Hairpinning

- "NAT UDP Unicast Requirements"
- F. Audet and C. Jennings (guest lecture last year)

More formally, a NAT that supports hairpinning forwards packets originating from an internal address, $X1:x1$, destined for an external address $X2':x2'$ that has an active mapping to an internal address $X2:x2$, back to that internal address, $X2:x2$. Note that typically $X1'$ is the same as $X2'$.

Furthermore, the NAT may present the hairpinned packet with either an internal ($X1:x1$) or an external ($X1':x1'$) source IP address and port. Therefore, the hairpinning NAT behavior can be either "External source IP address and port" or "Internal source IP address and port". "Internal source IP address and port" may cause problems by confusing implementations that expect an external IP address and port.

Example confusion

- Two nodes (A, B) share a switch behind a NAT
- A sends TCP SYN to an external address for B
- B responds with SYN+ACK
- What can happen if NAT did not rewrite A's SYN to have an external address and port? Will B's packet traverse the NAT?

Simple Traversal of UDP through NATs

STUN (RFC 3849)

- “Simple Traversal of User Datagram Protocol (UDP) Through Network Address Translators (NATs)”
- Enables a node to
 - Determine if it is behind a NAT, and if so, what kind
 - Obtain a public IP address/port pair
- Client-server protocol, requires no changes to NATs
- STUN server coordinates


STUN Binding Requests

- Node sends BR to STUN server
- STUN sends a response that has the address and port it sees
 - If different than node's local address and port, it's behind a NAT
- Node can probe to see what kind of NAT
 - Ask STUN server to respond from different address: no response, address restricted; different response, symmetric
 - Ask STUN server to respond from different port: no response, port restricted
- When does STUN not work?


NAT Hole-Punching

- STUN doesn't work when
 - A or B are behind symmetric NAT
 - A and B are behind same NAT without hairpinning
 - What about restricted cone, port-restricted NATs?
- Parallel approaches
 - A and B report their local IP address to server S_x
 - Server tells the other the address/port pair (L, G)
 - A tries to send UDP packets to (L_B, G_B) using L_A
 - B tries to send UDP packets to (L_A, G_A) using L_B


NAT Hole-Punching Example




NAT Hole-Punching Example




NAT Hole-Punching Example



NAT Hole-Punching Example




NAT Hole-Punching Example



Multiple NAT Layers

- Common for consumer Internet: ISP has internal NAT, end user places another NAT
- Requires hairpinning at ISP NAT



The New Hourglass

