

Administrivia

- **Midterm exam Thursday**
 - Open book, Open notes, no electronic devices allowed
 - Feel free to print out and bring lecture slides
- **SCPD students:**
 - Email `cs144-staff@scs.stanford.edu` with your exam monitor information
 - Please ensure the email subject is “exam monitor”
- **Any other students with special exam needs**
 - Please email `cs144-staff` to make arrangements


Outline

- DNS architecture
- DNS protocol and resource records (RRs)
- Record types: A, NS, glue, MX, SOA, CNAME
- Reverse lookup
- Load balancing
- DNS security

Parsing a URL



Motivation



- **Users can't remember IP addresses**
 - Need to map symbolic names (`www.stanford.edu`) → IP addr
- **Implemented by library functions & servers**
 - `getaddrinfo()` talks to server over UDP (sometimes TCP)
- **Actually, more generally, need to map symbolic names to values**

hosts.txt system

- **Originally, hosts were listed in a file, `hosts.txt`**
 - Email global network administrator when you add a host
 - Administrator mails out new `hosts.txt` file every few days
- **Would be completely impractical today**
 - `hosts.txt` today would be huge (Gigabytes)
 - What if two people wanted to add same name?
 - Who is authorized to change address of a name?
 - People need to change name mappings more often than every few days (e.g., Dynamic IP addresses)


Goals of DNS

- **Scalability**
 - Must handle huge number of records
 - Potentially *exponential* in name size—because custom software may synthesize names on-the-fly
- **Distributed control**
 - Let people control their own names
- **Fault-tolerance**
 - Old software assumed `hosts.txt` always there
 - Bad potential failure modes when name lookups fail
 - Minimize lookup failures in the face of other network problems

The good news

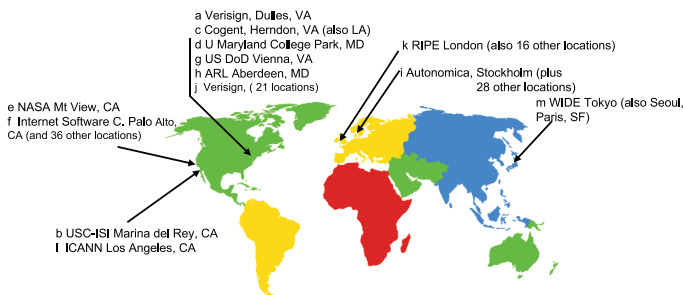
- **Properties that make DNS goals easier to achieve:**
 1. **Read-only or read-mostly database**
 - People typically look up hostnames much more often than they are updated
 2. **Loose consistency**
 - When adding a machine, may be okay if info takes minutes or hours to propagate
- **These suggest approach w. aggressive caching**
 - Once you have looked up hostname, remember result
 - Don't need to look it up again in near future

Domain Name System (DNS)




- **Break namespace into a bunch of zones**
 - . ("root"), edu., stanford.edu., cs.stanford.edu., ...
 - Zones separately administered \Rightarrow **delegation**
 - Parent zones tell you how to find servers for subdomains.
- **Each zone served from several replicated servers**

Root servers



- **Root (and TLD) servers must be widely replicated**
 - For some, use various tricks like IP anycast

DNS software architecture



- **Two types of query**
 - Recursive
 - Non-Recursive
- **Apps make recursive queries to local DNS server (1)**
- **Local server queries remote servers non-recursively (2, 4, 6)**
 - Aggressively caches result
 - E.g., only contact root on first query ending .umass.edu

DNS protocol

- **TCP/UDP port 53**
- **Most traffic uses UDP**
 - Lightweight protocol has 512 byte UDP message limit
 - retry w. TCP if UDP fails (e.g., reply truncated)
- **TCP requires message boundaries**
 - Prefix all messages w. 16-bit length
- **Bit in query determines if query is recursive**

Resource records

- **All DNS info represented as resource records (RR):**

name [TTL] [class] type rdata

 - *name* – domain name (e.g., www.stanford.edu.)
 - *TTL* – time to live in seconds
 - *class* – for extensibility, usually IN (1) "Internet"
 - *type* – type of the record
 - *rdata* – resource data dependent on the *type*
- **Two important DNS RR types:**
 - **A** – Internet address (IPv4)
 - **NS** – name server
- **Example resource records (dig stanford.edu):**

```
stanford.edu. 1800 IN A 171.67.216.14
stanford.edu. 1800 IN A 171.67.216.16
stanford.edu. 172800 IN NS Argus.stanford.edu.
...
```

Some implementation details

- How does local name server know root servers?

- Need to configure name server with *root cache* file
- Contains root name servers and their addresses

```
.          3600000  NS   A.ROOT-SERVERS.NET.
A.ROOT-SERVERS.NET. 3600000  A    198.41.0.4
.          3600000  NS   B.ROOT-SERVERS.NET.
B.ROOT-SERVERS.NET. 3600000  A    128.9.0.107
...
```

- How do you get addresses of other name servers

- To lookup names ending `.stanford.edu.`, ask `Argus.stanford.edu.`
- Chicken and egg problem:
How to get `Argus.stanford.edu.`'s address?
- Solution: **glue** records – A records in parent zone
- Name servers for `edu.` have A record of `Argus.stanford.edu.`

Glue Record Example

- Look up `www.scs.stanford.edu` assuming no cache

```
dig +norec www.scs.stanford.edu @a.root-servers.net
dig +norec www.scs.stanford.edu @a.edu-servers.net
dig +norec www.scs.stanford.edu @argus.stanford.edu
dig +norec www.scs.stanford.edu @ns1.fs.net
```

- Get intermediary results for `.edu`, `stanford.edu`, `scs.stanford.edu`, and `www.scs.stanford.edu`

- Where are the glue records?

Structure of a DNS message [RFC 1035]

Header	
Question	the question for the name server
Answer	RRs answering the question
Authority	RRs pointing toward an authority
Additional	RRs holding additional information

- Same message format for queries and replies
 - Query has zero RRs in Answer/Authority/Additional sections
 - Reply includes question, plus has RRs
- Authority allows for delegation
- Additional for glue + other RRs client might need

Header format

0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5
ID															
QR	Opcode		AA		TC	RD		RA	Z	RCODE					
QDCOUNT															
ANCOUNT															
NSCOUNT															
ARCOUNT															

- **QR** – 0=query, 1=response
- **OPCODE** - 0=standard query
- **RCODE** – error code
- **AA**=authoritative answer, **TC**=truncated, **RD**=recursion desired, **RA**=recursion available

Encoding of RRs

NAME															
TYPE															
CLASS															
TTL															
RDLENGTH															
RDATA															

Encoding of domain names

- A DNS name consists of a series of labels
 - `www.stanford.edu.` has 3 labels: `www`, `stanford`, and `edu`
 - Labels can contain letters, digits, and “-”, but should not start or end with “-”
 - Maximum length 63 characters
 - Encoded as length byte followed by label
 - Last label always empty (zero-length) label
- Names are case insensitive
 - But server must preserve case of question in replies
 - Example: request `www.sTANford.EDu`, look at authority

Name compression



- **Observation: many common suffixes in DNS messages**
 - Particularly because of case preservation rule
- **Allow pointer labels to re-use suffixes**
 - Recal label starts with length byte (0-63)
 - If value $\geq 0xc0$ (192), subtract $0xc000$ from first *two* bytes, and treat as pointer into message

Other Records

- **Start of Authority (SOA) record**
 - States administrative information for a zone
 - `dig stanford.edu soa`
 - Tells you how long you can cache negative results
- **Mail Exchange (MX) record**
 - For historical reasons, mail does not have to use A records directly
 - Example: `ping scs.stanford.edu`
 - No such host, but you can still mail CS144 staff there
 - `dig scs.stanford.edu mx`

Reverse Lookups

- Remember *traceroute...*
- Traceroute can learn names of hosts through *reverse lookup*
- `128.30.2.121 → 121.2.30.128.in-addr.arpa`
- PTR record points to canonical name
- **Example:**
 - `tinyos.stanford.edu → sing.stanford.edu`
 - `sing.stanford.edu → 171.67.76.65`
 - `65.76.67.171.in-addr.arpa → sing.stanford.edu`

Secondary servers

- **Availability requires geographically disparate replicas**
 - E.g., I ask MIT to serve `scs.stanford.edu`
- **Typical setup: One master many slave servers**
- **How often to sync up servers? Trade-off**
 - All the time \Rightarrow high overhead
 - Rarely \Rightarrow stale data
- **Put trade-off under domain owner's control**
 - Fields in SOA record control secondary's behavior
 - Primary can unilaterally change SOA
 - To speed propagation, primary can also notify secondary of change, providing a hint to refresh sooner [RFC 1996]

CNAME records

- **CNAME record specifies an alias:**
 - `name [TTL] [IN] CNAME canonical-name`
 - As if any RR's associated w. *canonical-name* also for *name*
 - Can look up with `AI_CANONNAME` flag to `getaddrinfo`
- **Examples, to save typing:**
 - `wb.scs.stanford.edu. CNAME williamsburg-bridge.scs.stanford.edu.`
 - `mb.scs.stanford.edu. CNAME manhattan-bridge.scs.stanford.edu.`
- **CNAME precludes any other RRs for name**
 - E.g., might want: `david.com CNAME david.stanford.edu`
 - Illegal, because `david.com` would need NS records
- **Note answer section can have CNAME for query name + other RR(s) for canonical-name**
 - But don't point MXes to CNAMEs, as no A recs in additional section (try `bad-mx.scs.stanford.edu.`)

Mapping addresses to names

- **PTR records specify names**
 - `name [TTL] [IN] PTR "ptrdname"`
 - *name* – somehow encode address...how?
 - *ptrdname* – domain name for this address
- **IPv4 addrs stored under in-addr.arpa domain**
 - Reverse name, append `in-addr.arpa`
 - To look up `171.66.3.9 → 9.3.66.171.in-addr.arpa.`
 - Why reversed? Delegation!
- **IPv6 under ip6.arpa**
 - Historical note: ARPA funded original Internet
 - Acronym now re-purposed [RFC 3172]:
Address and Routing Parameter Area

2-minute stretch



SRV records

- Service location records
 - `_service._proto.name [...]` SRV *prio weight port target*
 - `_service` – E.g., `_sip` for SIP (VOIP) protocol
 - `_proto` – `_tcp` or `_udp`
 - `name` – domain name record applies to
 - `prio` – as with MX records, lower # → higher priority
 - `weight` – within priority, affects randomization of order
 - `port` – TCP or UDP port number (particularly useful for SIP)
 - `target` – Server name, for which client needs A record
- Like a generalization of MX records for arbitrary services



- SPF is based on envelope sender address
 - Nice because available earlier in SMTP protocol
 - So some users can reject forged mail while some accept
- Microsoft proposed competing standard, *Sender ID* [RFC 4406]
 - Instead of simple language, used XML monstrosity
 - Instead of envelope sender, extract address from message
- No agreement between camps, couldn't standardize
 - Compromise: kill XML, but use address in message
 - But Microsoft patented extracting address from message!

Using DNS for load-balancing

- Can have multiple RR of most types for one name
 - Required for NS records (for availability)
 - Useful for A records
 - (Not legal for CNAME records)
- Servers rotate order in which records returned
 - `getaddrinfo` returns a linked list of `addrinfo` structures
 - Most apps just use first address returned
 - Even if your name server caches results, clients will be spread amongst servers
- Example: `dig cnn.com` multiple times

TXT records

- Can place arbitrary text in DNS
 - `name [TTL] [IN] TXT "text" ...`
 - `text` – whatever you want it to mean
- Great for prototyping new services
 - Don't need to change DNS infrastructure
- Example: `dig gmail.com txt`
 - What's this? SPF = "sender policy framework" (previously known as "sender permitted from")
 - Much spam is forged email
 - SPF specifies IP addresses allowed to send mail from `@gmail.com`
 - Can have incremental deployment
 - Only mail servers must change, DNS can stay the same
 - Now SPF standardized (sort of), has RR type 99 [RFC 4408]

SPF vs. Sender ID (continued)

- Compromise 2: Have two competing standards
 - After a few years, see which standard more widely used
- Use different formats for SPF vs. Sender ID
 - Start SPF records with string `"v=spf1 "`
 - Start Sender ID records with string `"spf2.0/pra "`
- SPF had a head start—lots of sites had adopted it
- Dirty trick appeared in final draft of Sender ID
 - If no `spf2.0/pra` record present, but see `v=spf1`, treat `v=spf1` as if it were a sender ID record
 - Causes sender ID machines to reject mail from SPF sites (E.g., if you use SPF and post to mailing list, some recipients will reject)
 - Thwarts idea of independent experiment

DNS redirection for content distribution

- Play with akamai and `www.microsoft.com`

Classless in-addr delegation

- How to delegate on non-byte boundary?
- Solution: Use CNAME records
 - So-called *classless* in-addr delegation
- Example:

```
1.3.66.171.in-addr.arpa. CNAME 1.ptr.your-domain.com.  
2.3.66.171.in-addr.arpa. CNAME 2.ptr.your-domain.com.  
3.3.66.171.in-addr.arpa. CNAME 3.ptr.your-domain.com.
```


DNS exploits

- July 29, 2008, **Bruce Schneier**:
[Despite the best efforts of the security community, the details of a critical internet vulnerability discovered by Dan Kaminsky about six months ago have leaked.](#)
- One of the basic problems: DNS caching
 - If you can poison the cache, the damage stays
 - Who knows how far it spreads...

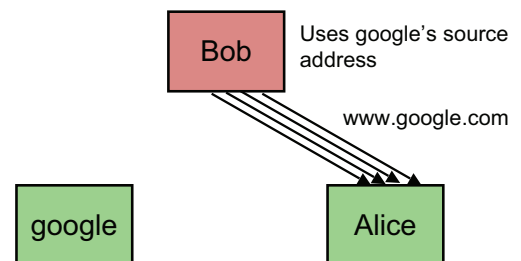
DNS exploit example

- Alice wants to look up `www.google.com`
- Bob the attacker knows
- Bob knows source address/port, destination address/port
- Bob generates a spoof response: `www.google.com` is `www.evil.com`
- Challenge: Bob has to guess **Query ID**
- If Bob guesses, RR can stay in Alice's cache a long time

Exploit Example



Exploit Example




Countermeasures

- Choose good QIDs (used to be incremented, now randomly generated), 16 bits
- Randomize source port, 16 bits
- Some protection, but only makes it take longer, networks are faster each day

Another exploit

- DNS clients used to trust all responses
- **Problem: glue records and helpful A records**
 - Ask NS of `evil.com` for `www.evil.com`
 - Says `www.evil.com` is a CNAME for `www.amazon.com`
 - Provides A record for `www.amazon.com`

Exploit Example



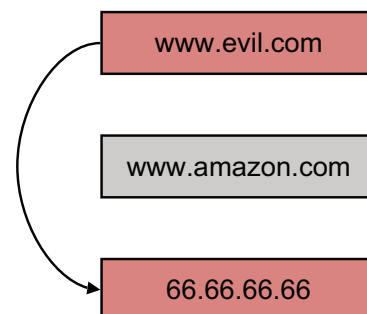
It gets worse

- Glue records can overwrite standard A records
- Even if you have a good A record for `www.amazon.com`, it's overwritten
- E.g., Server wants name of my IP address
 - Looks up `66.66.66.66.in-addr.arpa`
- I say nameserver for `66.66.66.66.in-addr.arpa` is `www.amazon.com`
 - Include glue A record for `www.amazon.com` in my reply

Solution 1

- Only use glue records for duration of query
 - Cache only end-to-end traversal of pointers, not intermediate steps
- In CNAME example `www.evil.com` will point to evil server
 - `www.amazon.com` will not point to evil server
- In `in-addr.arpa` example, can lie about hostname
 - But I can lie anyway
 - Have to check reverse lookup result by doing forward lookup

Example



Solution 2: bailiwick checking

- Only pay attention to answers for the domain you've asked
- Response from `evil.com` can't tell you the A record for `google.com`
- Ask `google.com` for `www.google.com`
- Opponent can still race, but at least it's not deterministic

Kaminsky exploit

- Make winning the race easier
- Brute force attack
- Force Alice to look up `AAAA.google.com`, `AAAB.google.com`, etc.
- Forge CNAME responses for each lookup, inserting A record for `www.google.com`
- Circumvents bailiwick checking

Solution: signatures

- Signature: cryptographic way to prove a party is who they say they are (more later in quarter)
- Requires a chain of trust
- Whom do you trust to sign DNS?
- DNSSEC extensions may finally be deployed soon [\[RFC 4033\]](#)

DNS Overview

- Distributed system for mapping names to values (e.g., IP addresses)
- Read-dominated workload allows caching
- Name structure allows distribution, independent administration
- Caching means bad data can stay a long time
- Standard protocol does not authenticate response is from server: DNSSec does